

Lab 2 Writeup - CSE490R

Tarkan Al-Kazily, Yuta Araki, Dilraj Devgun

May 2019

1 Introduction

In this lab, we implemented three different kinds of motion controllers: a PID controller, a Pure Pursuit controller, and a Model Predictive controller. This writeup goes over some of the choices we made in tuning these controllers, and compares their performance.

Along with this writeup we have a large store of videos in the `videos` folder. Here are a few of the ones you should watch.

Video	Description
<code>sim_pid.mp4</code>	Simulation of our PID controller.
<code>sim_pp.mp4</code>	Simulation of our Pure Pursuit controller.
<code>sim_mpc.mp4</code>	Simulation of our MPC controller.
<code>purepursuit_real.mp4</code>	Our Pure Pursuit controller on the real robot.
<code>mpc_tuned.mp4</code>	Our MPC controller on the real robot.
<code>mpc_tuned_avoid_wall.mp4</code>	Our MPC controller avoiding real obstacles.

2 PID Controller

The first controller we implemented was a PID controller, which stands for Proportional Integral Derivative controller. To reduce complexity for this lab, we were instructed not to implement the integral gain, making our controller more of a PD controller.

The initially provided values for K_p and K_d were effective before tuning. Its performance curve is given in Figure 1. In tuning, we prioritized the controller quickly reaching the set point of the path (given by the cross track error going to zero) as well as minimizing the overshoot of the car. We first determined the K_p gain until our controller would reach the reference path in the desired time interval (approximately 100 to 150 ms on the plot), ignoring the consequence of overshooting, and then increased the K_d gain until the controller no longer overshoot the set-point. Figure 2 shows the performance of our controller with only K_p non-zero (tuned to 0.3) and Figure 3 shows the performance of the final, tuned controller with both $K_p = 0.4$ and $K_d = 0.3$. We did end up increasing the K_p after tuning K_d to improve the time it takes the controller to reach the setpoint.

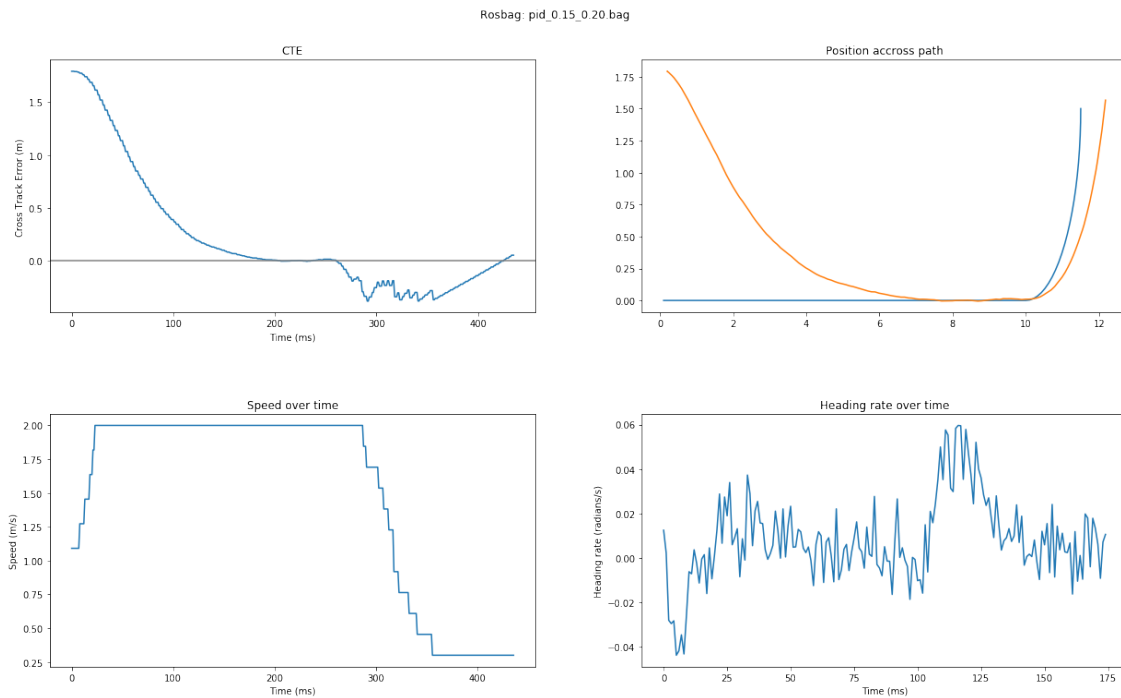


Figure 1: Original PID performance

We found that increasing K_p had the effect of taking sharper turns towards the set-point, up to a point where the minimum turning radius of the car was achieved. Increasing K_d would cause the car to straighten out nearer to the reference path which would reduce overshooting.

As mentioned, our final K_p and K_d values are $K_p = 0.4$ and $K_d = 0.3$.

To implement the derivative computation of our controller, we used the analytic approach with $\sin(\psi_e)$. In terms of the controller's performance it does very well at ensuring that as the error decreases and the car reaches the reference path that the car is pointing in the direction of the reference path, and won't cross it rapidly. We chose to implement this first because of its simplicity over numerical differentiation (as it doesn't rely on tracking previous error values). It also has the benefit that it isn't as impacted by noise in sampling the error, since using numerical integration if the samples for the derivative aren't smooth unexpectedly large derivative terms can negatively impact performance. However, due to time constraints we did not have time to implement and test numerical derivation methods.

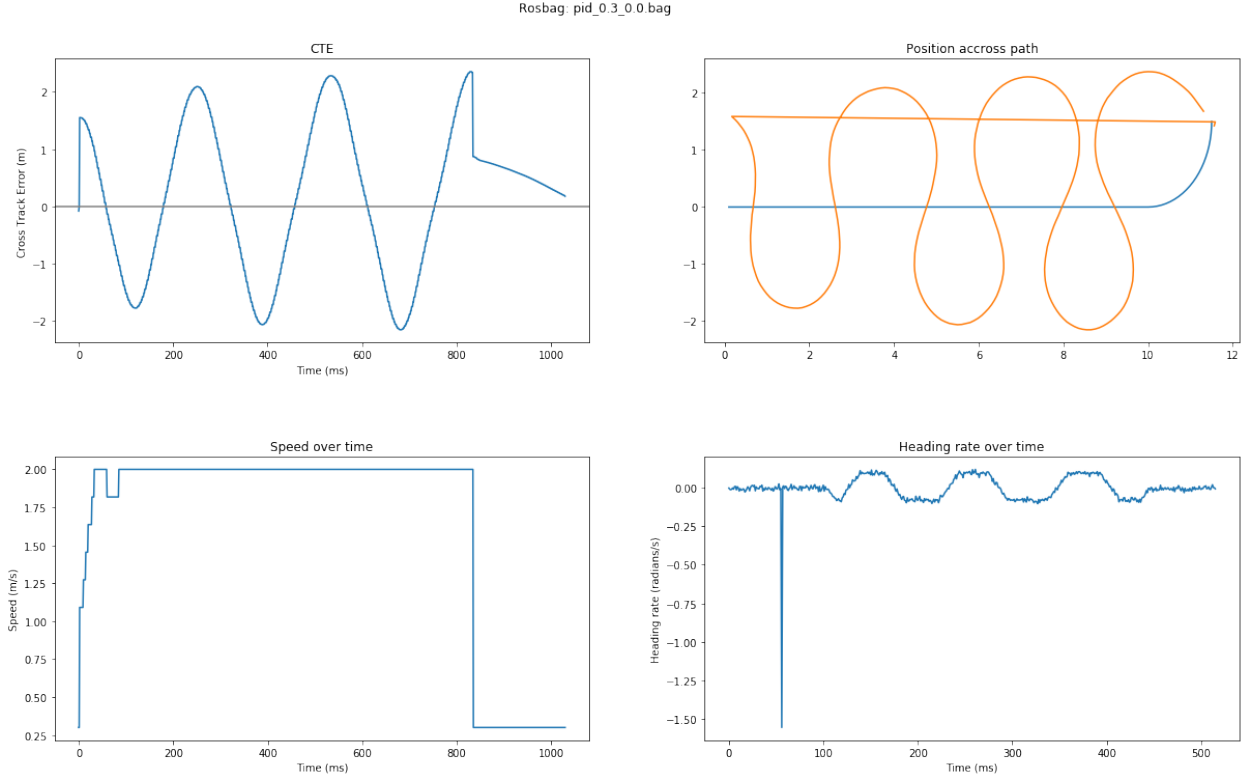


Figure 2: Approximate desired K_p performance (0.3) with $K_d = 0.0$

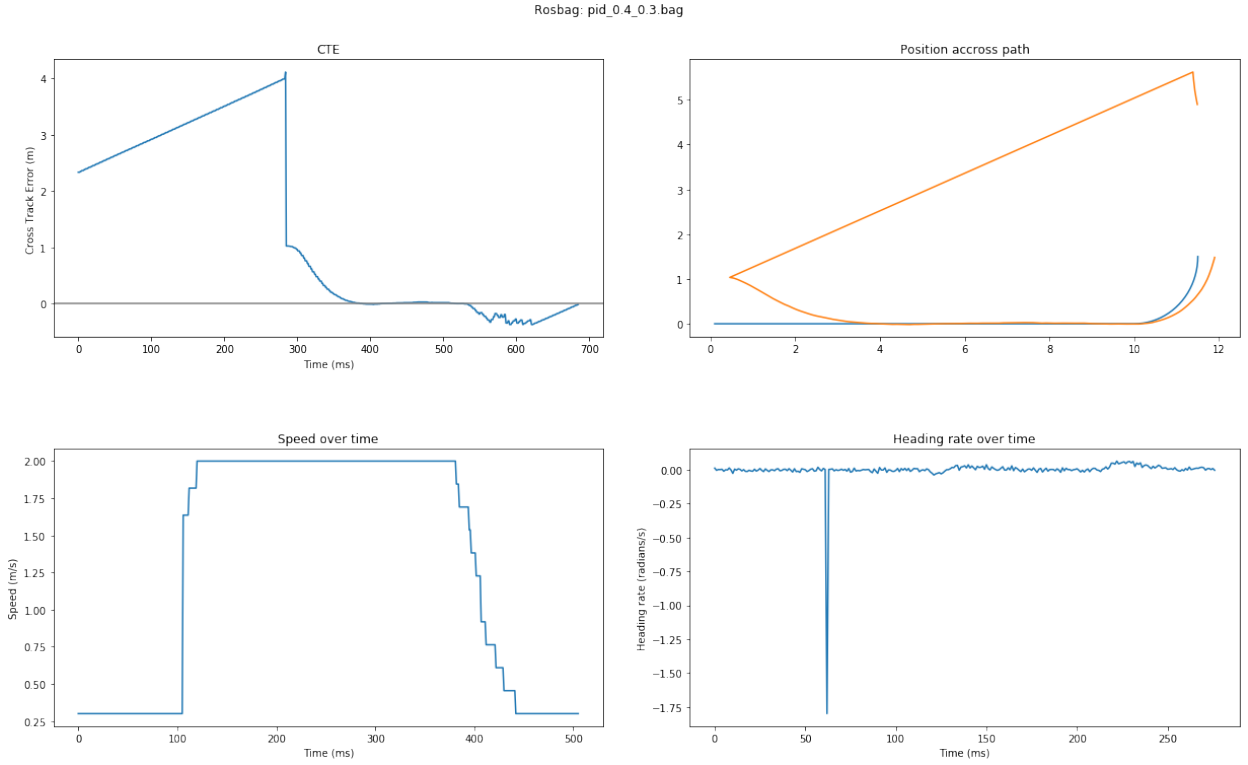


Figure 3: Tuned K_p and K_d performance ($K_p = 0.4, K_d = 0.3$)

3 Pure Pursuit Controller

We started with vary short look-ahead distance of 0.1 meters (the performance for which is given in Figure 4), which did not work too well as it forced the car to take super sharp turns repeatedly and caused car to move in a circular path. From there, we kept increasing the look-ahead distance until controller was able to get car back on track in reasonable manner, which was at 0.9 (see Figure 5). We tried higher values for look-ahead, which caused car to cut corners (see Figure 6 to see the performance of a lookahead of 1.5 meters). Look-ahead distance less than 0.7 or greater than 1.2 did not perform well due to reasons stated above.

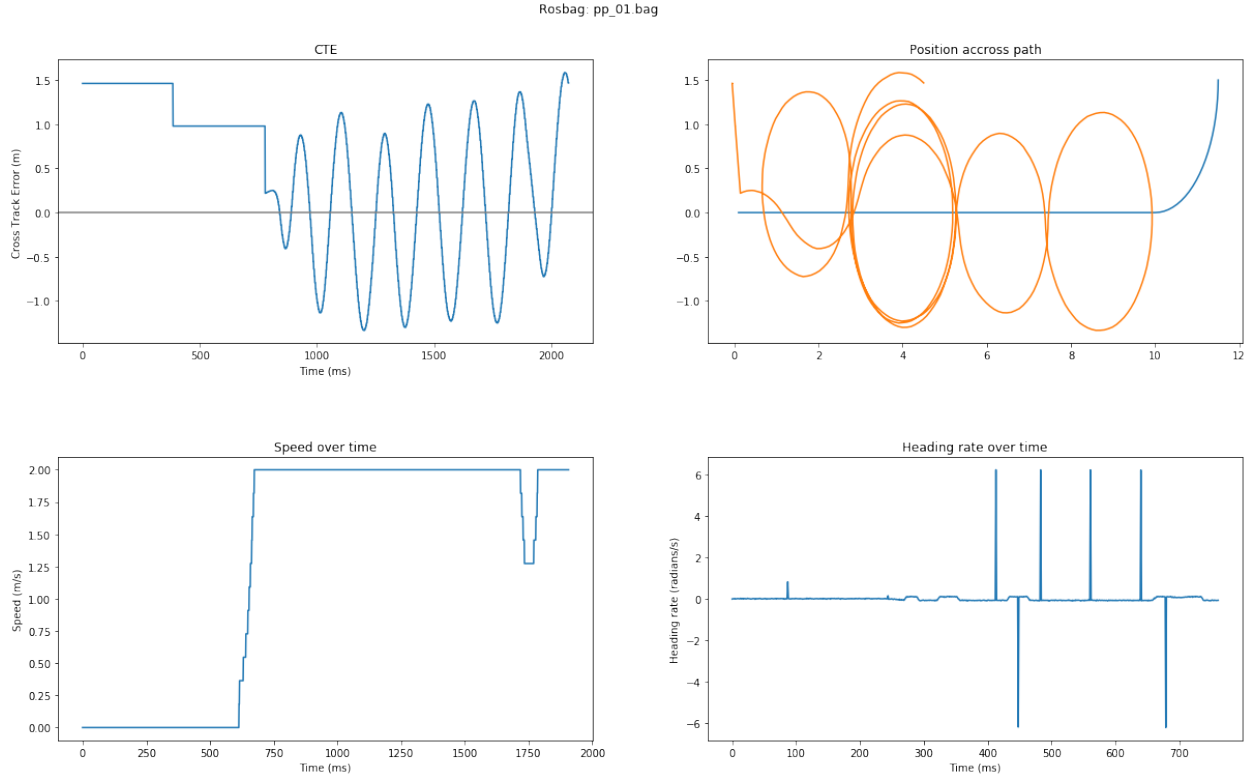


Figure 4: Pure Pursuit performance with lookahead of 0.1 meters

Varying turning radius of the circle path did not have significant affect on the performance of our controller, as long as turn radius on the path was larger than car's capability. Testing on various turning radius did not necessary help narrowing down the optimal look-ahead distance, however, it did help us confirm that look-ahead distance we chose was good enough.

Varying the `desired_speed` parameter did not affect the performance of our controller. Faster speed made the car to overshoot a little but there was no significant overshoot. Lower speed made decreased the overshooting.

When running the pure pursuit controller on the real robot, we found that our existing lookahead value worked well and didn't need additional tuning.

Comparing PD controller and Pure Pursuit controller, it seems like PD controller is more generalizable and Pure Pursuit controller performs better on constant curvature (e.g. circle path).

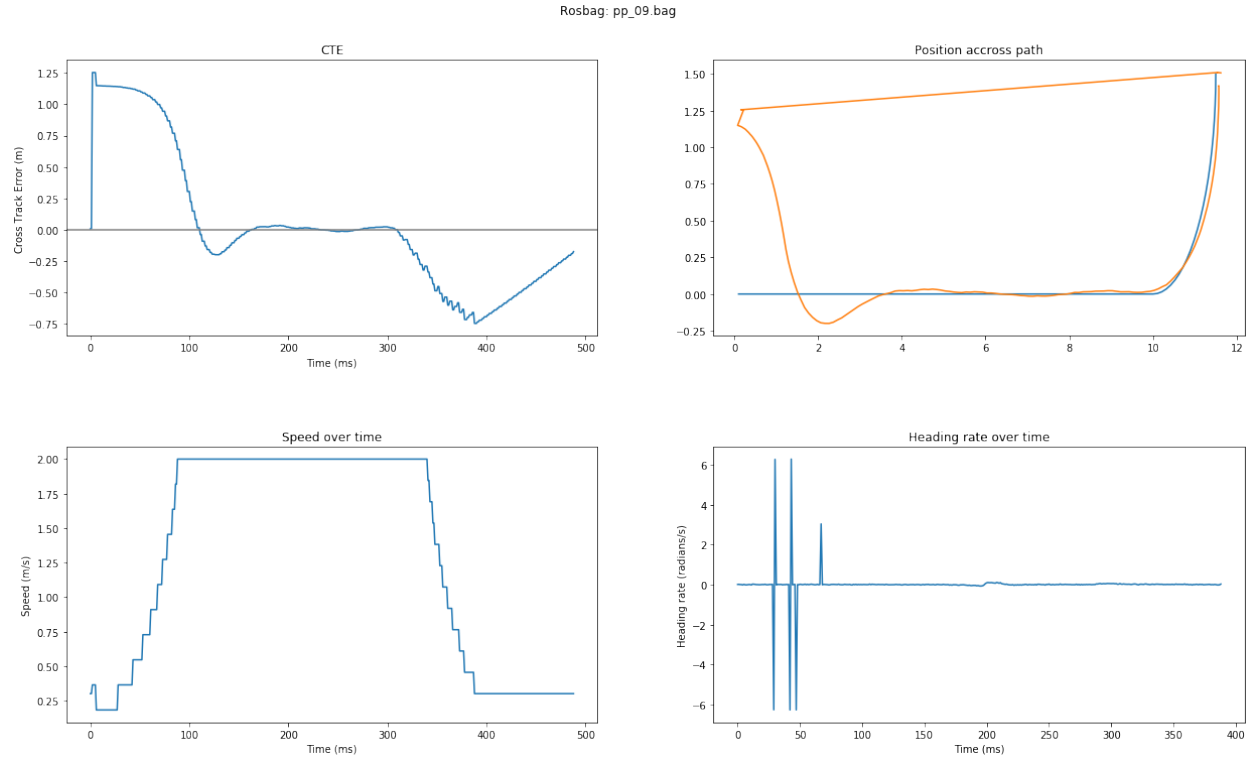


Figure 5: Final Pure Pursuit performance with lookahead of 0.9 meters

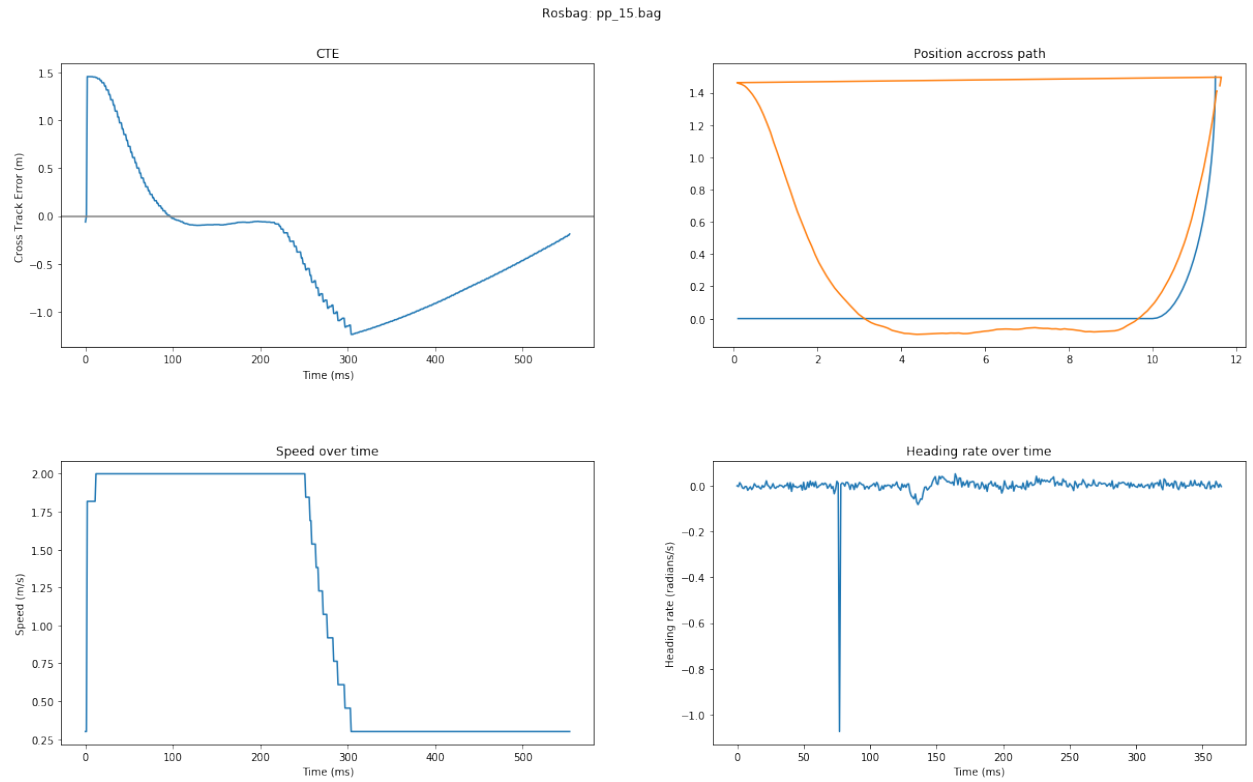


Figure 6: Pure Pursuit performance with lookahead of 1.5 meters

4 Model Predictive Controller

A plot of our trajectories and cost function heatmap is given in Figure 7. Initially we used constant curvature trajectories with minimal variation between paths due to the simplicity of development, but decided to implement a variation where after a certain percentage of our actions the steering angle is set to straight. This change further penalizes paths that put the car directly in front of an obstacle in the short term because it doesn't assume that the car will continue to turn as part of the future path planning.

We began tuning our MPC by testing on the initial parameters for K and T which caused our robot not to avoid obstacles since we were not looking far enough in the future to adjust our planned motion. From this point we doubled our T by twice to 16 and also again to 32. The results for these values were much better, yielding results that we expected where the robot would navigate around forthcoming objects. In addition, we also ran our simulation using a roll out of 100 which caused our robot to search too far into the future which would make us stop early and make too harsh of adjustments to what is immediately impending for the robot to navigate. The performances of these tests in simulation are given in Figures 8, 9, 10, and 11 respectively.




Figure 7: Our MPC Trajectories and Heatmap

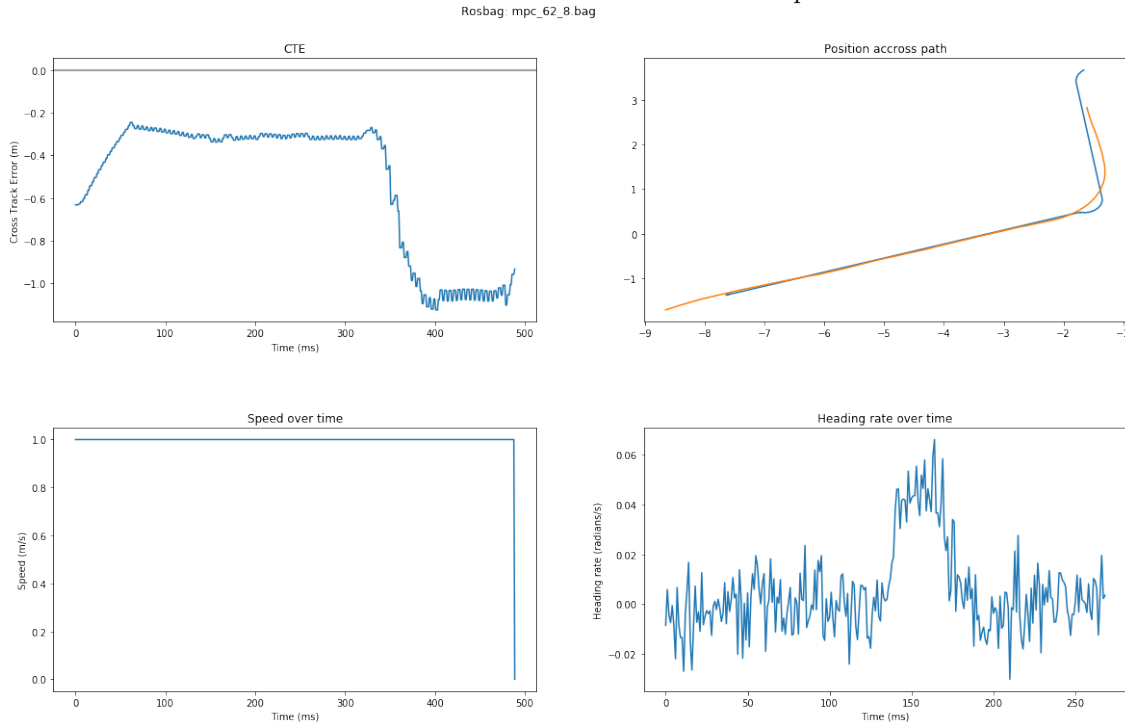


Figure 8: Simulated K and T performance ($K = 62, T = 8$)




Figure 9: Simulated K and T performance ($K = 62, T = 16$)

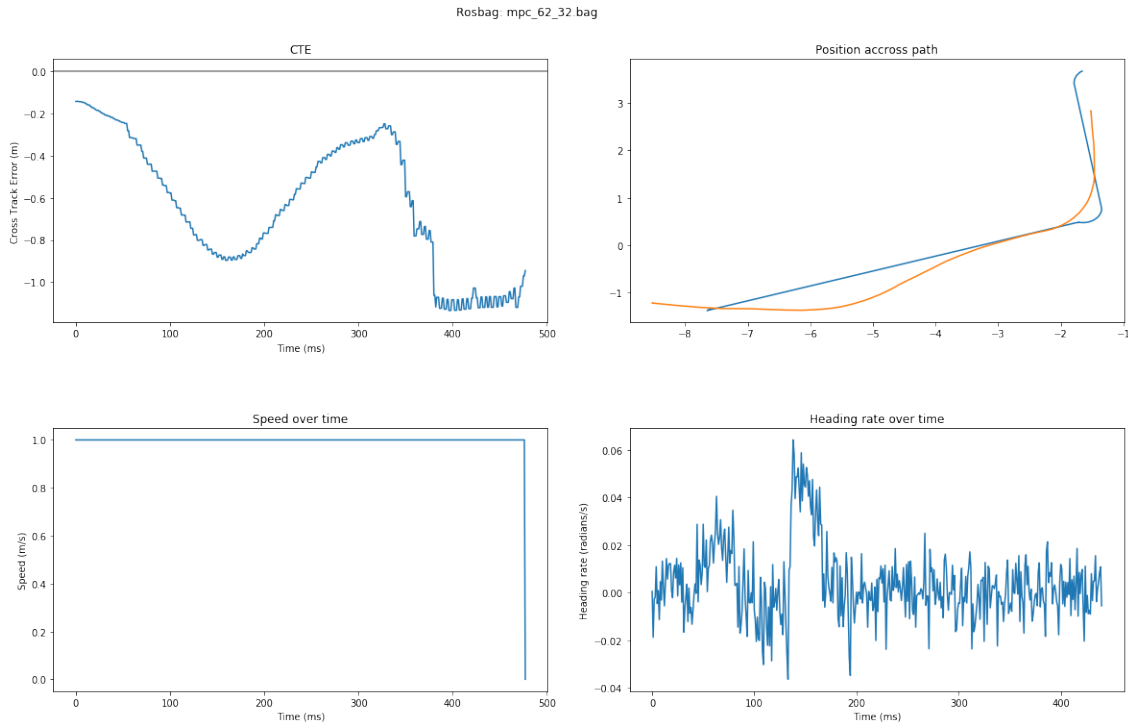


Figure 10: Simulated K and T performance ($K = 62, T = 32$)

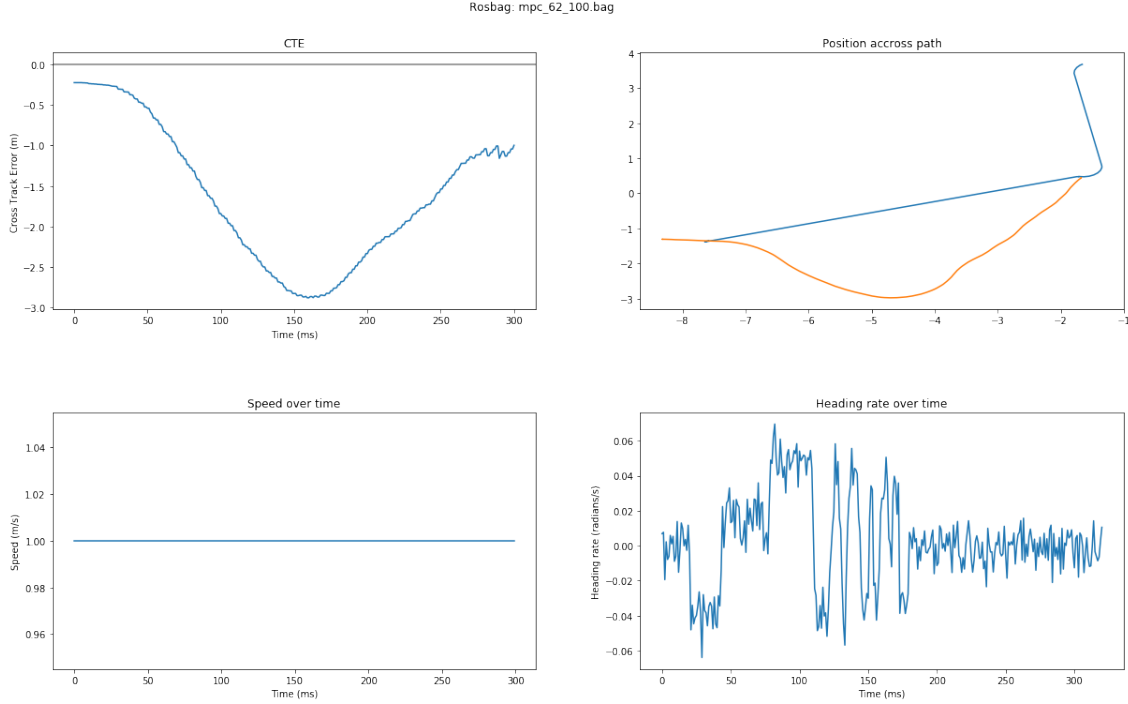


Figure 11: Simulated K and T performance ($K = 62, T = 100$)

Once we had a strong initial value for T , we tested it on the real robot to avoid the cabinet in CSE 022. We found that with the default cost function having such a high T value (32) was problematic with the existing dt parameter (which was in timesteps of .1 second). We experimented with smaller T values, but as part of tuning on the real robot found that the dt value was better in smaller timesteps, finally choosing dt to be 0.4 seconds. With this we increased T to have the similar rollout distance as before. We also increased K to guarantee a greater density of poses further from the current pose. Our final T and K values are 60 and 150 respectively.

Finally, we experimented slightly with our cost function, with a focus on tweaking the collision weight so that it would reliably avoid obstacles in the real. Initially we tried increasing the collision weight by a factor of 10, and found that resulted in the robot avoiding the path too much. A weight of $5 \cdot 10^5$ was more appropriate when equally distributed over all the collision poses. We finally implemented a scaled weight distribution that more heavily penalizes collisions for rollouts fewer timesteps from the current pose. This had a very positive effect on our car's avoidance performance, and reduced the case where paths that extended beyond the other side of the obstacle were not penalized enough. It did require tuning our scaling factor properly, along with the overall collision weight, which we did by trying to increase these values until the car would more reliably avoid obstacles.

For the lookahead principle, we used the same method as in our PID controller, in which we found our next waypoint by finding the closest waypoint and then returning the waypoint that is the lookahead distance in front of that point. This is similar in concept to Pure Pursuit's lookahead principle, but uses the lookahead distance from the reference pose closest to our current pose, as opposed to the distance from the current pose.

5 Overall Findings

The PID controller was best on straightaways, but did adequately on turns with small curvature, since its K_p gain would be enough to make it turn consistently. We believe that if we added an integral term to the PID controller that it would more closely follow a circular reference path, since over time the integral term would be able to correct for the negative feedback of the derivative factor.

The Pure Pursuit controller worked best on large, consistent turns such as the circle reference path, which makes sense due to the controller’s inherent assumption that all motion the car makes travels in a circle. It would often oscillate on straight-aways more than other controllers, but could operate faster than the other controllers.

The MPC controller stood out as the most generalizable controller among our choices, and did well regardless of the scenario. We believe this is due to the generalizability of optimizing the cost of multiple paths as part of planning each control action. It has the obvious strength that it can adapt its path in the presence of minor obstacles on the path. It was the most robust to failures because of this, easily returning to the path if it over shot or avoided an obstacle.

All our controllers performed well in the sim-to-real transition, but there was definitely some discrepancy. Our PID controller was one of the weakest, because the real car’s response behaves differently from the simulation, and therefore requires more tuning to overcome factors not present in the simulation. An integral term would also help here, as we noticed that the settle point of the PID controller was not directly on the reference path. The pure pursuit transitioned to the real world the best. Finally, our MPC controller took a lot of tuning in the real world to choose more effective K and T values for the real world case.